

University of Edinburgh, School of Mathematics

Incomplete Data Analysis, 2020/2021

Assignment 2 – Solutions

Ted Ladas

Github Repo: https://github.com/TedOiler/ida_assignment_2

Question 1

$F(y; \theta) = 1 - e^{-y^2/(2\theta)}$ and $y \geq 0, \theta \geq 0$ and \

$$X_i = \begin{cases} Y_i & Y_i \leq C \\ C & Y_i > C \end{cases}$$

$$R_i = \begin{cases} 1 & Y_i \leq C \\ 0 & Y_i > C \end{cases}$$

a) Calculation of $\hat{\theta}_{MLE}$

We need to calculate the pdf of the data given theta, the survival function of the upper censor point, and the Likelihood of the data given theta. In order to reach the conclusion of the exercise I will define first these.

$$X_i = Y_i R_i + C(1 - R_i) \iff X_i^2 = Y_i^2 R_i^2 + 2Y_i R_i C(1 - R_i) + C^2(1 - R_i)^2$$

We know that $R_i = 0 \vee R_i = 1$ therefore $2Y_i R_i C(1 - R_i) = 0$

If $R_i = 0$ then $R_i^2 = R_i$ and $(1 - R_i)^2 = (1 - R_i)$

If $R_i = 1$ then $R_i^2 = R_i$ and $(1 - R_i)^2 = (1 - R_i)$

Therefore $R_i^2 = R_i$ and $(1 - R_i)^2 = (1 - R_i)$

So, Finally: $X_i^2 = Y_i^2 R_i + C^2(1 - R_i)$

We also have to have an expression for the Survival function:

$$S(C; \theta) = 1 - F(y; \theta) \iff S(C; \theta) = e^{-y^2/(2\theta)}$$

And finally we need to derive the pdf of the data

$$f(y; \theta) = \frac{\partial}{\partial y} F(y; \theta) \iff f(y; \theta) = \frac{1}{\theta} y e^{-y^2/(2\theta)}$$

So we can start:

$$L(\theta) = \prod_i \{ (f(y; \theta))^{r_i} (S(C; \theta))^{(1-r_i)} \} = \prod_i \{ (\frac{1}{\theta} y e^{-y^2/(2\theta)})^{r_i} (e^{-C^2/(2\theta)})^{(1-r_i)} \} =$$

$$\theta^{-\sum r_i} \sum y_i^{r_i} e^{-\sum y_i^2 r_i + C^2(1-r_i)} = \theta^{-\sum r_i} \sum y_i^{r_i} e^{-\sum X_i^2}$$

$$\ln L(\theta) = \ell(\theta) = -\sum r_i \ln \theta - \frac{1}{2\theta} \sum X_i^2 + \ln \sum y_i^{r_i}$$

$$\frac{\partial}{\partial \theta} \ell(\theta) = 0 \iff -\frac{\sum r_i}{\theta} + \frac{\sum X_i^2}{2\theta^2} + 0 = 0 \iff \frac{-2\theta \sum r_i + \sum X_i^2}{2\theta^2} = 0$$

and since $\theta \geq 0$

$$-2\theta \sum r_i + \sum X_i^2 = 0 \iff \theta_{MLE} = \frac{\sum X_i^2}{2 \sum r_i}$$

b) We know that $I(\theta) = -\mathbb{E} \left[\frac{d^2 \ell(\theta)}{d\theta^2} \right]$

and also that $\mathbb{E}[R_i] = 1P(R_i = 1) + 0P(R_i = 0) = P(Y \leq C) = F(C; \theta) = 1 - e^{-C^2/(2\theta)}$ and therefore $\mathbb{E}[1 - R_i] = 1 - \mathbb{E}[R_i] = e^{-C^2/(2\theta)}$

so if we decompose $\sum X_i^2$ again we have:

$$\begin{aligned} I(\theta) &= -\frac{\sum \mathbb{E}[r_i]}{\theta^2} + \frac{\sum \mathbb{E}[Y_i^2 r_i]}{\theta^3} + \frac{\sum \mathbb{E}[C^2(1-r_i)]}{\theta^3} = \\ &= n/\theta^2(1 - e^{-C^2/(2\theta)}) + n/\theta^3(-C^2 e^{-C^2/(2\theta)} + 2\theta(1 - e^{-C^2/(2\theta)}) - e^{-C^2/(2\theta)}) \\ I(\theta) &= \frac{n}{\theta^2}(1 - e^{-C^2/(2\theta)}) \end{aligned}$$

c) Since we know that $\theta_{MLE} \sim N(\theta, I(\theta)^{-1})$ we can Normalize θ_{MLE} and produce the 95% confidence interval as follows

$$Z = \frac{\theta_{MLE} - \theta}{\sqrt{I(\theta)^{-1}}} \sim N(0, 1)$$

$$P(z_{a/2} \leq Z \leq z_{a/2}) = 1 - a$$

$$a = 0.05 \text{ hence } z_{-a/2} = -1.959964, z_{a/2} = 1.959964$$

$$[z_{-a/2} \sqrt{I(\theta)^{-1}}, z_{a/2} \sqrt{I(\theta)^{-1}}]$$

Question 2

a)

Like question 1a our likelihood is the product of the censored and the noncensored data. Since we know that the Normal distribution belongs to the exponential family, we could recompute the likelihood and prove the statement. However, this involves a lot of calculations, mostly to calculate the CDF of a normal distribution and to then manipulate it by multiplying it with the PDF in order to derive a distribution of X_i . Fortunately, there is an easier way to show this result.

We know that $Y_i \sim N(\mu, \sigma^2)$ If we can show that $X_i \sim N(\mu, \sigma^2)$ then the result is trivial.

$$X_i = Y_i R_i + D(1 - R_i)$$

Since our X is an affine transformation of a normal, we can write it as $X = g(Y)$ where g linear. Secondly we notice that the data are right censored which means: $P(Y < D; \mu, \sigma^2) = 1 - S(Y; \mu) = F(Y)$ Together with the fact that R_i have a binary nature of $0 \vee 1$ we have our result.

So. $X_i \sim N(\mu, \sigma^2)$

Hence: $\log L(\mu, \sigma^2 | x, r) = \log \prod [\phi(x_i; \mu, \sigma^2)^{r_i} \times \Phi(x_i; \mu, \sigma^2)] = \sum_{i=1}^n \{r_i \log \phi(x_i, \mu, \sigma^2) + (1 - r_i) \Phi(x_i; \mu, \sigma^2)\}$

b) $\rightarrow \mu = 5.559766$

```
# https://cran.r-project.org/web/packages/maxLik/maxLik.pdf
load(file = "dataex2.Rdata")

ll <- function(m, df) {
  mu = m
  sigma = sqrt(1.5)
  x_i = df[, 1]
  r_i = df[, 2]
  likelihood_f = sum(r_i * dnorm(x_i, mean = mu, sd = sigma, log = TRUE) + (1 -
    r_i) * pnorm(x_i, mean = mu, sd = sigma, log = TRUE)) #log=TRUE: to avoid underflow, instead of
  # Nice to have: modify the optimization newton-raphson function we wrote on
  # statistical programming to optimize this problem as well.
  return(likelihood_f)
}

start = c(mu = 1)
mle = maxLik(logLik = ll, df = dataex2, start = start, method = "nr") # Newton-Raphson method is the fastest
summary(mle)

## -----
## Maximum Likelihood estimation
## Newton-Raphson maximisation, 4 iterations
## Return code 1: gradient close to zero
## Log-Likelihood: -339.7295
## 1 free parameters
## Estimates:
##      Estimate Std. error t value Pr(> |t|)
## mu  5.559777    0.08761    63.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Question 3

Based on the lecture ‘Likelihood based inference with incomplete data’ we know that a missing data mechanism is ignorable for likelihood inference iff:

- 1) The missing data are MAR (or MCAR)
- 2) the parameter ψ which is the parameter of the missing mechanism and the parameters of the model θ are distinct.

On examples (a) and (b) we notice that point (2) is true because. So we only have to consider the first point to answer the question. More detail for (c) on the corresponding subsection.

What we are doing on each example is the following We express the probability of a datapoint missing as a linear model of various other variables and parameters. The `logit()` function serves the purpose of transforming the support of our probability $R_f \in [0, 1]$ to $R_{logitf} \in (-\infty, +\infty)$ so that our line makes sense.

(a) Missing data mechanism \rightarrow MAR.

This is true because the missingness depends on Y_1 which is fully observed. Therefore the mechanism is ignorable for likelihood-based estimation

(b) Missing data mechanism \rightarrow MNAR

This is true because the missingness depends on Y_2 which is not fully observed. Therefore the mechanism is **not** ignorable for likelihood-based estimation

(c) Missing data mechanism \rightarrow MAR.

This is true because the missingness depends on Y_1 . However we notice that μ_1 is part of the missingness mechanism so condition (2) is not true. Therefore the mechanism is **not** ignorable for likelihood-based estimation

Question 4

Theoretical Implementation:

since we know that Y_i are **individual** draws from a Bernoulli Distribution we can calculate the likelihood as follows:

$$L(\beta) = \prod_{i=1}^N [p_i(\beta)^{y_i} (1 - p_i(\beta))^{1-y_i}]$$

where

$$\beta = [\beta_0 \ \beta_1]^\top$$

$$p_i(\beta) = \sigma(\beta_0 + \beta_1 x_i)$$

Therefore:

$$\ell(\beta) = \log L(\beta) = \sum_{i=1}^N [y_i \log(\frac{e^{\beta_0 + \beta_1 x_i}}{1 + e^{\beta_0 + \beta_1 x_i}}) + (1 - y_i) \log(\frac{1}{1 + e^{\beta_0 + \beta_1 x_i}})] = \sum_{i=1}^N [y_i (e^{\beta_0 + \beta_1 x_i}) - y_i \log(1 + (e^{\beta_0 + \beta_1 x_i})) - \log(1 + (e^{\beta_0 + \beta_1 x_i})) + y_i \log((1 + e^{\beta_0 + \beta_1 x_i}))] \iff$$

$$\ell(\beta) = \sum_{i=1}^N y_i (\beta^\top x_i) - \log(1 + e^{\beta^\top x_i})$$

So our Q function, keeping in mind that the missing data are from $m+1$ to n , without loss of generality:

$$\ell(\beta) = \sum_{i=1}^M y_i (\beta^\top x_i) + \sum_{i=m+1}^N \mathbb{E}[y_i] (\beta^\top x_i) - \log(1 + e^{\beta^\top x_i})$$

and since we know that $Y_i \sim \text{Bernoulli}(p_i(\beta))$ we know that

$$\mathbb{E}[y_i] = p_i(\beta)$$

$$\ell(\beta) = \sum_{i=1}^M y_i (\beta^\top x_i) + \sum_{i=m+1}^N p_i(\beta) (\beta^\top x_i) - \log(1 + e^{\beta^\top x_i})$$

Them we can maximize the Q function as in Question 2 with the maxLik function

```
load("dataex4.Rdata")

# modifying dataset to be easier to work with
df4 <- dataex4 %>% mutate(R = (Y == 0 | Y == 1) * 1) %>% replace_na(list(R = 0)) %>%
  replace_na(list(Y = 999))

# helper function to evaluate the sigmoid easier
sigmoid <- function(bb, xx) {
  b0 <- bb[1]
  b1 <- bb[2]
  sig = exp(b0 + b1 * xx) / (1 + exp(b0 + b1 * xx))
  return(sig)
}

# E step
q_fun <- function(bb, df) {
  b0 <- bb[1]
  b1 <- bb[2]
  xx <- df[, 1]
  yy <- df[, 2]
  rr <- df[, 3]
  q <- sum((rr * (yy * (b0 + b1 * xx)) - log(1 + exp(b0 + b1 * xx))) + (1 - rr) *
    (b0 + b1 * xx) * sigmoid(beta_old, xx))
  return(q)
}
```

```

# M step
epsilon = 1e-08 # tolerance set to machine precision
beta_old = c(1, -1)

repeat {
  beta_new = coef(maxLik(logLik = q_fun, df = df4, start = beta_old))
  convergence = sum(abs(beta_old - beta_new)) < epsilon
  if (convergence) {
    break
  }

  beta_old = beta_new
}

cat(beta_new)

```

```
## 0.9755261 -2.480384
```

Question 5

(a)

- Our mixture model is:

$$P(Y = y) = f(y) = pf_{\log Normal}(y; \mu, \sigma^2) + (1 - p)f_{exp}(y; \lambda)$$

$$Z_i = \begin{cases} 1 & \text{if } y_i \text{ comes from the logNormal} \\ 0 & \text{if } y_i \text{ comes from the Exponential} \end{cases}$$

- When we write θ we refer to the parameters of the mode which are:

$$\theta = (p, \mu, \sigma^2, \lambda)$$

- Thus the Likelihood function of the parameters with regards to our data and our latent variable is:

$$L(\theta; y, z) = \prod_{i=1}^N \left\{ \left[p \frac{1}{y\sqrt{2\pi\sigma^2}} e^{-\frac{(\log y_i - \mu)^2}{2\sigma^2}} \right]^{z_i} [(1 - p)\lambda e^{-\lambda y_i}] \right\}$$

- And so our log Likelihood $\ell(\theta; y, z)$ is:

$$\ell(\theta; y, z) = \log L(\theta; y, z) = \sum_{i=1}^N [z_i \{ \log p - \log y_i - \log \sqrt{2\pi\sigma^2} - \frac{(\log y_i - \mu)^2}{2\sigma^2} \}] + \sum_{i=1}^N [(1 - z_i) \{ \log(1 - p) + \log \lambda - \lambda y_i \}]$$

- So our $Q(\theta|\theta^{(t)})$ function for the E-step is:

$$Q(\theta|\theta^{(t)}) = \sum_{i=1}^N [\mathbb{E}[z_i|y_i, \theta^{(t)}] \{ \log p - \log y_i - \log \sqrt{2\pi\sigma^2} - \frac{(\log y_i - \mu)^2}{2\sigma^2} \}] + \sum_{i=1}^N [(1 - \mathbb{E}[z_i|y_i, \theta^{(t)}]) \{ \log(1 - p) + \log \lambda - \lambda y_i \}]$$

- We know that $\mathbb{E}[z_i|y_i, \theta^{(t)}] = P(z_i = 1|y_i, \theta^{(t)})$ and therefore we can apply Bayes theorem and the Law of total probability to obtain:

$$\mathbb{E}[z_i|y_i, \theta^{(t)}] = \frac{p^{(t)} \frac{1}{y_i \sqrt{2\pi\sigma^{2(t)}}} e^{-\frac{(\log y_i - \mu^{(t)})^2}{2\sigma^{2(t)}}}}{p^{(t)} \frac{1}{y_i \sqrt{2\pi\sigma^{2(t)}}} e^{-\frac{(\log y_i - \mu^{(t)})^2}{2\sigma^{2(t)}}} + (1 - p^{(t)}) \lambda^{(t)} e^{-\lambda^{(t)} y_i}} = \tilde{p}_i^{(t)}$$

- So:

$$Q(\theta|\theta^{(t)}) = \sum_{i=1}^N [\tilde{p}_i^{(t)} \{ \log p - \log y_i - \log \sqrt{2\pi\sigma^2} - \frac{(\log y_i - \mu)^2}{2\sigma^2} \}] + \sum_{i=1}^N [(1 - \tilde{p}_i^{(t)}) \{ \log(1 - p) + \log \lambda - \lambda y_i \}]$$

- Finally for the M-step we just need to differentiate the above Q expression with respect to the parameters θ and set the equal to zero, to obtain our $\theta^{(t+1)}$

$$\frac{\partial Q(\theta|\theta^{(t)})}{\partial p} = 0 \iff 1/p \sum_{i=1}^N \tilde{p}_i^{(t)} - 1/(1 - p) \sum_{i=1}^N (1 - \tilde{p}_i^{(t)}) = 0$$

$$p^{(t+1)} = 1/n \sum_{i=1}^N \tilde{p}_i^{(t)}$$

$$\frac{\partial Q(\theta|\theta^{(t)})}{\partial \mu} = 0 \iff \sum_{i=1}^N \tilde{p}_i^{(t)} \left[\frac{\log y_i - \mu}{\sigma^2} \right] = 0$$

$$\mu^{(t+1)} = \frac{\sum_{i=1}^N \tilde{p}_i^{(t)} \log y_i}{\sum_{i=1}^N \tilde{p}_i^{(t)}}$$

$$\frac{\partial Q(\theta|\theta^{(t)})}{\partial \sigma^2} = 0 \iff \sum_{i=1}^N \tilde{p}_i^{(t)} \left[\frac{(\log y_i - \mu)^2}{2\sigma^4} - \frac{1}{2\sigma^2} \right] = 0$$

$$\sigma^{2(t+1)} = \frac{\sum_{i=1}^N \tilde{p}_i^{(t)} (\log y_i - \mu^{(t+1)})^2}{\sum_{i=1}^N \tilde{p}_i^{(t)}}$$

$$\frac{\partial Q(\theta|\theta^{(t)})}{\partial \lambda} = 0 \iff \sum_{i=1}^N (1 - \tilde{p}_i^{(t)}) \left[\frac{1}{\lambda} - y_i \right] = 0$$

$$\lambda^{(t+1)} = \frac{\sum_{i=1}^N (1 - \tilde{p}_i^{(t)})}{\sum_{i=1}^N (1 - \tilde{p}_i^{(t)}) y_i}$$

(b)

We can easily take the expressions for the E and M step and derive the algorithm.

```
load("dataex5.Rdata")
mixture_model <- function(y, theta0, eps) {

  theta <- theta0
  p <- theta[1]
  mu <- theta[2]
  sigma_sq <- theta[3]
  lambda <- theta[4]
  std <- sqrt(sigma_sq)
  diff <- 1
  n <- length(y)

  repeat {
    theta_old <- theta0

    # E-step
    dist1 <- p * dlnorm(y, mu, std)
    dist2 <- (1 - p) * dexp(y, lambda)
    p_til <- dist1 / (dist1 + dist2)

    # M-step
    p_t <- mean(p_til)
    mu_t <- sum(p_til * log(y)) / sum(p_til)
    sigma_sq_t <- sum(p_til * (log(y) - mu_t)^2) / sum(p_til)
    lambda_t <- (n - sum(p_til)) / sum((1 - p_til) * y)

    theta <- c(p_t, mu_t, sigma_sq_t, lambda_t)
    diff <- sum(abs(theta - theta_old))
    if (diff > eps) {
      return(theta)
    }
  }
}
```

```
theta_start = c(0.1, 1, 0.5^2, 2)
res <- mixture_model(y = dataex5, theta0 = theta_start, eps = epsilon)
cat(res)
```

```
## 0.5092109 1.981153 0.8008751 1.256405
```

```
p <- as.double(res[1])
mu <- as.double(res[2])
sigma_sq <- as.double(res[3])
lambda <- as.double(res[4])
std <- sqrt(sigma_sq)
std <- sqrt(sigma_sq)

dist1 <- p * dlnorm(dataex5, mu, std)
dist2 <- (1 - p) * dexp(dataex5, lambda)
orange = rgb(1, 0.65, 0, 0.35)
blue = rgb(0.65, 0.75, 1)

hist(dataex5, xlab = "data", freq = FALSE, col = orange, ylim = c(0, 0.15), )
curve(p * dlnorm(x, mu, std) + (1 - p) * dexp(x, lambda), add = TRUE, col = blue,
      lwd = 4, type = "o")
legend(100, 0.13, legend = c("Histogram", "MM pdf"), fill = c(orange, blue), bty = "o")
```

Histogram of dataex5

